

```
In [1]: import pandas as pd
import warnings
warnings.filterwarnings('ignore')
import yfinance as yf
import matplotlib.pyplot as plt
```

```
In [2]: # Extraemos Los dos DataSets
```

```
ibm_walmart = yf.download(tickers = 'IBM WMT', period = '5y', interval = '1mo', rounding = True)
ibm_walmart.head()
```

```
[*****100%*****] 2 of 2 completed
```

```
Out[2]:
```

Price	Adj Close		Close		High		Low		Open		Volume	
Ticker	IBM	WMT	IBM	WMT	IBM	WMT	IBM	WMT	IBM	WMT	IBM	WMT
Date												
2019-11-01 00:00:00+00:00	101.47	36.75	128.54	39.70	133.02	41.79	126.91	39.03	128.59	39.31	67751825	349861800
2019-12-01 00:00:00+00:00	102.35	36.67	128.15	39.61	130.42	40.71	124.94	39.14	128.54	39.72	75159179	295268100
2020-01-01 00:00:00+00:00	109.75	35.48	137.41	38.16	139.38	39.96	127.34	37.56	129.06	39.62	118459187	383722500
2020-02-01 00:00:00+00:00	99.38	33.37	124.43	35.89	151.77	39.98	120.80	34.79	137.91	38.30	123581342	397242300
2020-03-01 00:00:00+00:00	85.59	35.22	106.05	37.87	130.11	42.69	86.58	34.00	125.00	35.87	198819809	956864400

```
In [3]: ibm_walmart = ibm_walmart['Close']
ibm_walmart.head()
```

```
Out[3]:
```

Ticker	IBM	WMT
Date		
2019-11-01 00:00:00+00:00	128.54	39.70
2019-12-01 00:00:00+00:00	128.15	39.61
2020-01-01 00:00:00+00:00	137.41	38.16
2020-02-01 00:00:00+00:00	124.43	35.89
2020-03-01 00:00:00+00:00	106.05	37.87


```
In [34]: rendimientos = IBM.pct_change()
rendimientos = rendimientos.dropna()
rendimientos
```

```
Out[34]:
```

Ticker	IBM
Date	
2019-11-01 00:00:00+00:00	0.013453
2019-11-04 00:00:00+00:00	0.015822
2019-11-05 00:00:00+00:00	0.001596
2019-11-06 00:00:00+00:00	0.006448

257 rows × 4 columns

```
In [4]: # graficamos
ibm_walmart['IBM'].plot()
ibm_walmart['WMT'].plot()
plt.legend(['IBM', 'WMT'])
plt.title('Comparación de precios IBM - Walmart')
plt.show()
```



```
In [6]: # Rendimientos
rend_ibm_wmt = ibm_walmart.pct_change()
rend_ibm_wmt
```

Out[6]:

Ticker	IBM	WMT
Date		
2019-11-01 00:00:00+00:00	NaN	NaN
2019-12-01 00:00:00+00:00	-0.003034	-0.002267
2020-01-01 00:00:00+00:00	0.072259	-0.036607
2020-02-01 00:00:00+00:00	-0.094462	-0.059486

```
In [7]: # eliminamos Los NAN

rend_ibm_wmt = rend_ibm_wmt.dropna()
rend_ibm_wmt
```

```
Out[7]:
```

Ticker	IBM	WMT
Date		
2019-12-01 00:00:00+00:00	-0.003034	-0.002267
2020-01-01 00:00:00+00:00	0.072259	-0.038807
2020-02-01 00:00:00+00:00	-0.094482	-0.059488
2020-03-01 00:00:00+00:00	-0.147714	0.055189
2020-04-01 00:00:00+00:00	0.131919	0.089978


### Aplicamos la Nomenclatura y componentes de las series de tiempo

```
In [8]: # Calculamos la correlación

correlacion = rend_ibm_wmt['IBM'].corr(rend_ibm_wmt['WMT'])
print('Correlación entre IBM y Walmart: ', correlacion)

Correlación entre IBM y Walmart: 0.2770463660898834
```

```
In [9]: # graficamos
plt.scatter(rend_ibm_wmt['IBM'], rend_ibm_wmt['WMT'])
plt.show()
```



- ¿Existe alguna correlación entre los precios de las acciones de ambas empresas?
- Observando la Gráfica y la correlación, vemos que no hay relación entre los precios de ambas acciones

```
In [11]: # podemos corroborar esta correlación con el modelo de regresión simple
import statsmodels.api as sm
rend_ibm_wmt = sm.add_constant(rend_ibm_wmt)
rend_ibm_wmt
```

```
Out[11]:
```

	const	IBM	WMT
Date			
2019-12-01 00:00:00+00:00	1.0	-0.003034	-0.002267
2020-01-01 00:00:00+00:00	1.0	0.072259	-0.038607
2020-02-01 00:00:00+00:00	1.0	-0.094482	-0.059488
2020-03-01 00:00:00+00:00	1.0	-0.147714	0.055189

```
In [12]: regresion = sm.OLS(rend_ibm_wmt['IBM'], rend_ibm_wmt[['const', 'WMT']]).fit()
```

```
In [13]: print(regresion.summary())
```

```

OLS Regression Results
=====
Dep. Variable:          IBM    R-squared:                0.077
Model:                  OLS    Adj. R-squared:           0.061
Method:                 Least Squares    F-statistic:          4.739
Date:                  Wed, 30 Oct 2024    Prob (F-statistic):    0.0336
Time:                  21:40:01    Log-Likelihood:        77.323
No. Observations:      59    AIC:                  -150.6
Df Residuals:          57    BIC:                  -146.5
Df Model:              1
Covariance Type:       nonrobust
=====
                    coef    std err          t      P>|t|      [0.025    0.975]
-----
const             0.0054      0.009      0.610      0.544     -0.012     0.023
WMT               0.3491      0.160      2.177      0.034      0.028     0.670
=====
Omnibus:            0.243    Durbin-Watson:          2.003
Prob(Omnibus):      0.886    Jarque-Bera (JB):        0.198
Skew:               -0.131    Prob(JB):                0.906
Kurtosis:           2.892    Cond. No.                18.6
=====
```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

```
In [14]: # calculamos el índice de correlación
autocorrelacion = rend_ibm_wmt['IBM'].autocorr()
autocorrelacion
```

```
Out[14]: -0.014759191403178598
```

```
In [19]: autocorrelacion = ibm_walmart['IBM'].autocorr()  
autocorrelacion
```

```
Out[19]: 0.9267057150281207
```

```
In [20]: IBM = yf.download(tickers = 'IBM', period = '5y', interval = '1d', rounding = True)  
IBM = IBM['Close']
```

```
[*****100%*****] 1 of 1 completed
```

```
In [22]: # aplicamos Las diferencias consecutivas  
IBM_diff = IBM.diff()  
IBM_diff
```

```
Out[22]:
```

Ticker	IBM
Date	
2019-10-31 00:00:00+00:00	NaN
2019-11-01 00:00:00+00:00	1.72
2019-11-04 00:00:00+00:00	2.05
2019-11-05 00:00:00+00:00	2.04

```
In [24]: autocorr_diff = IBM_diff['IBM'].autocorr()  
autocorr_diff
```


```
Out[24]: -0.0291691800424556
```

```
In [26]: # generacion de La funcion de utocorrelacion  
from statsmodels.tsa.stattools import acf  
from statsmodels.graphics.tsaplots import plot_acf
```

```
In [27]: acf_array = acf(IBM)  
acf_array
```

```
Out[27]: array([1.          , 0.99345266, 0.98670695, 0.97950292, 0.97211519,  
                0.96441491, 0.95518912, 0.9463187 , 0.93675462, 0.92763549,  
                0.91782713, 0.9077393 , 0.89782   , 0.88793764, 0.87833055,  
                0.86865654, 0.85938242, 0.85033942, 0.84149532, 0.8327977 ,  
                0.82426098, 0.81624398, 0.80811862, 0.80021112, 0.79196781,  
                0.78383544, 0.77621711, 0.76869266, 0.76093226, 0.75337665,  
                0.74611517])
```


```
In [28]: plot_acf(IBM, alpha = 0.05)
plt.show()
```



```
In [35]: rendimientos = rendimientos.dropna()
acf_rend = acf(rendimientos)
print(acf_rend)
```

```
[ 1.00000000e+00 -6.88977982e-02  5.26747789e-02 -2.36832375e-02
 -2.45647724e-02  7.65622144e-02 -7.96770841e-02  1.34745538e-01
 -1.06537573e-01  9.48817480e-02  2.36408333e-05 -1.31366271e-02
 -1.90885706e-02 -4.21350835e-02  4.16877608e-02 -9.06925715e-02
  5.64904688e-02 -3.00199264e-02 -1.45474234e-02 -4.64945311e-03
 -4.55443796e-02  2.22553267e-02 -6.55883319e-02  4.64812027e-02
 -4.35328914e-02 -5.05222615e-02 -2.08037798e-02  3.28265723e-02
 -7.18823228e-03  1.74159733e-02 -3.66072399e-02]
```

```
In [36]: plot_acf(rendimientos, alpha = 0.05, lags = 20)
plt.show()
```



```
In [58]: # Prueba de caminata aleatoria, hay 2 hipotesis
# H0: Se tiene un proceso de Caminata Aleatoria
# Ha: No se tiene un proceso de Caminata Aleatoria.

from statsmodels.tsa.stattools import adfuller
import numpy as np


# generación de 500 pasos aleatorios

np.random.seed(1)
steps = np.random.normal(loc = 0.001, scale = 0.01, size = 200) +1
```

```
In [95]: steps[0] = 2
p = 100 * np.cumprod(steps)
```

```
In [96]: # Graficamos Los precios simulados

plt.plot(p, color = 'green')
plt.title('Proceso de Caminata Aleatoria')
plt.show()
```



```
In [97]: resultado = adfuller(p)
print('El resultado de p en la prueba es :', resultado[1])

El resultado de p en la prueba es : 0.9989986444709306
```


**0.99, indica que no hay suficiente evidencia para rechazar la hipótesis nula de que la serie no es estacionaria.**

```
In [98]: # Prueba Dickey Fuller para los precios de IBM
resultado = adfuller(IBM)
print('El resultado de p en la prueba es:', resultado[1])


El resultado de p en la prueba es: 0.7347657589182868
```

```
In [99]: rolling_ibm = IBM.rolling(30).mean()

plt.figure(figsize=(8,6))
plt.plot(IBM, label = 'Precios de IBM', color = 'green')
plt.plot(rolling_ibm, label = 'Media Movil', color = 'red')
plt.title('Precios Historicos de IBM vs Promedios Moviles a 30 días')
plt.legend()
plt.show()
```




```
In [100]: # Descomposición de La serie de tiempo
IBM.plot(figsize=(10,7)).grid()
```





```
In [101]: decomposed_ibm = sm.tsa.seasonal_decompose(IBM, period = 30)
```

```
In [102]: figure = decomposed_ibm.plot()  
plt.grid()  
plt.show()
```



## Realizamos Predicciones

\*\* Predicción Simple mediante promedios móviles

```
In [103]: IBM.shape
```

```
Out[103]: (1258, 1)
```

```
In [104]: len_train = int(1259 * 0.8)
len_train
```

```
Out[104]: 1007
```

```
In [105]: # Grupos de entrenamiento y prueba
train = IBM[0:len_train]
test = IBM[len_train:]
```

```
In [106]: len(train)
```

```
Out[106]: 1007
```

```
In [107]: len(test)
```

```
Out[107]: 251
```

```
In [137]: #y_pred = IBM.copy()
#y_pred = pd.DataFrame(y_pred)
ibm_data = IBM.copy()
ibm_data = pd.DataFrame(ibm_data)
```

```
In [138]: ibm_data['promedio_movil'] = ibm_data.rolling(30).mean()
```

```
In [112]: y_pred = ibm_data.copy()
y_pred['pronostico'][29] = 'NaN'
for contador in range(30, 1258):
    y_pred['pronostico'][contador] = y_pred['pronostico'][contador]
y_pred[0:35]
```


```
Out[112]:
```

```
In [113]: y_pred = pd.DataFrame(y_pred['pronostico'][1007:])
y_pred
```

```
Out[113]:
```

pronostico	
Date	
2023-11-01 00:00:00+00:00	141.566
2023-11-02 00:00:00+00:00	141.553667
2023-11-03 00:00:00+00:00	141.586667
2023-11-06 00:00:00+00:00	141.669667
2023-11-07 00:00:00+00:00	141.856

```
In [114]: # Graficamos
plt.figure(figsize = (16,7))
plt.grid()
plt.plot(train, label = 'train')
plt.plot(test, label = 'test')
plt.plot(y_pred['pronostico'], label = 'Predicción', color = 'green')
plt.legend()
plt.title('Pronosticos mediante Medias Moviles')
plt.show()
```



```
In [115]: test = pd.DataFrame(test)
```

```
In [116]: res = pd.concat([test.reset_index(drop=True), y_pred.reset_index(drop=True)], axis=1)
res.columns = ['Observado', 'pronostico']
```

```
In [118]: # Evaluamos Los pronosticos

from sklearn.metrics import mean_squared_error
```

```
In [119]: # primero obtenemos el rmse (raiz cuadrado medio, mean_squared_error)

rmse = np.sqrt(mean_squared_error(test, y_pred)).round(2)
mape = np.round(np.mean(np.abs((test['IBM'] - y_pred['pronostico']) / test['IBM']))) * 100, 2)
```

```
In [120]: print('RMSE: ', rmse, 'Mape: ', mape)
```

RMSE: 9.58 Mape: 4.34

**En promedio cuando hacemos pronosticos, nos equivocamos 9.58 dls y eso equivale a 4.34% que es nuestro error promedio absoluto hacia arriba o hacia abajo**

```
In [134]: # Pronóstico: tomar el último valor calculado del promedio móvil como estimación para el siguiente día
ultimo_promedio_movil = ibm_data['promedio_movil'].iloc[-1]

# Crear una fecha adicional para el 1 de noviembre de 2024
proximo_dia = pd.Timestamp('2024-11-01')

# Crear un nuevo DataFrame para almacenar la fecha y el pronóstico
pronostico = pd.DataFrame({'Close': [np.nan], 'promedio_movil': [ultimo_promedio_movil]}, index=[proximo_dia])

# Concatenar el pronóstico al final de los datos existentes
ibm_data = pd.concat([ibm_data, pronostico])

# Imprimir el pronóstico
print("Pronóstico para el 1 de noviembre de 2024:", ultimo_promedio_movil)

Pronóstico para el 1 de noviembre de 2024: 224.29266666666666
```

in